

Animal Classification with Keras

Mehdi ABOUZAID
Damien TOOMEY
Weihao ZHOU

INSA – National Institute of Applied Sciences

April 24, 2019

Plan

- 1 Introduction
- 2 CNN
- 3 Xception
- 4 VGG-16
- 5 ResNet-50
- 6 Conclusion
- 7 References

Plan

1 Introduction

2 CNN

3 Xception

4 VGG-16

5 ResNet-50

6 Conclusion

7 References

Introduction

Some information



- Article:

Animal Recognition and Identification with Deep Convolutional Neural Networks for Automated Wildlife Monitoring

- Studied models:

	ILSVRC	Parameters	Trainable Layers	Validation Set	
				Top-1 Accuracy	Top-5 Accuracy
AlexNet	2012	62,378,344	8	0.633	0.846
VGG-16	2014	138,357,544	16	0.713	0.901
ResNet-50	2015	25,636,712	50	0.749	0.921
Xception	2017	22,910,480	71	0.790	0.945




- Vocabulary:

kernel = filter = receptive field = mask

Introduction

Our dataset



Input	Output	Number of Images
	butterfly	1991
	cow	2039
	squirrel	2013

Plan

- 1 Introduction
- 2 **CNN**
- 3 Xception
- 4 VGG-16
- 5 ResNet-50
- 6 Conclusion
- 7 References

CNN: VGG-16 as example

Input



$$\mathcal{I} = \begin{pmatrix} I_{1,1} & I_{1,2} & I_{1,3} & I_{1,4} \\ I_{2,1} & I_{2,2} & I_{2,3} & I_{2,4} \\ I_{3,1} & I_{3,2} & I_{3,3} & I_{3,4} \\ I_{4,1} & I_{4,2} & I_{4,3} & I_{4,4} \end{pmatrix}$$

$$\overbrace{\begin{pmatrix} R_{1,1} & R_{1,2} & R_{1,3} & R_{1,4} \\ R_{2,1} & R_{2,2} & R_{2,3} & R_{2,4} \\ R_{3,1} & R_{3,2} & R_{3,3} & R_{3,4} \\ R_{4,1} & R_{4,2} & R_{4,3} & R_{4,4} \end{pmatrix} \quad \begin{pmatrix} G_{1,1} & G_{1,2} & G_{1,3} & G_{1,4} \\ G_{2,1} & G_{2,2} & G_{2,3} & G_{2,4} \\ G_{3,1} & G_{3,2} & G_{3,3} & G_{3,4} \\ G_{4,1} & G_{4,2} & G_{4,3} & G_{4,4} \end{pmatrix} \quad \begin{pmatrix} B_{1,1} & B_{1,2} & B_{1,3} & B_{1,4} \\ B_{2,1} & B_{2,2} & B_{2,3} & B_{2,4} \\ B_{3,1} & B_{3,2} & B_{3,3} & B_{3,4} \\ B_{4,1} & B_{4,2} & B_{4,3} & B_{4,4} \end{pmatrix}}$$

\mathcal{I}_{red}

\mathcal{I}_{green}

\mathcal{I}_{blue}

CNN: VGG-16 as example

2D Convolution : only 3x3 kernel, stride 1, zero padding of thickness 1



\mathcal{I}_{red}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & R_{1,1} & R_{1,2} & R_{1,3} & R_{1,4} & 0 \\ 0 & R_{2,1} & R_{2,2} & R_{2,3} & R_{2,4} & 0 \\ 0 & R_{3,1} & R_{3,2} & R_{3,3} & R_{3,4} & 0 \\ 0 & R_{4,1} & R_{4,2} & R_{4,3} & R_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

$Kernel[:, :, 0]$

$$\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix}$$

\mathcal{I}_{green}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & G_{1,1} & G_{1,2} & G_{1,3} & G_{1,4} & 0 \\ 0 & G_{2,1} & G_{2,2} & G_{2,3} & G_{2,4} & 0 \\ 0 & G_{3,1} & G_{3,2} & G_{3,3} & G_{3,4} & 0 \\ 0 & G_{4,1} & G_{4,2} & G_{4,3} & G_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

$Kernel[:, :, 1]$

$$\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}$$

\mathcal{I}_{blue}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & B_{1,1} & B_{1,2} & B_{1,3} & B_{1,4} & 0 \\ 0 & B_{2,1} & B_{2,2} & B_{2,3} & B_{2,4} & 0 \\ 0 & B_{3,1} & B_{3,2} & B_{3,3} & B_{3,4} & 0 \\ 0 & B_{4,1} & B_{4,2} & B_{4,3} & B_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

$Kernel[:, :, 2]$

$$\begin{bmatrix} w_{19} & w_{20} & w_{21} \\ w_{22} & w_{23} & w_{24} \\ w_{25} & w_{26} & w_{27} \end{bmatrix}$$

Goal : learn the weights in the kernels

CNN: VGG-16 as example



2D Convolution : only 3x3 kernel, stride 1, zero padding of thickness 1

\mathcal{I}_{red}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & R_{1,1} & R_{1,2} & R_{1,3} & R_{1,4} & 0 \\ 0 & R_{2,1} & R_{2,2} & R_{2,3} & R_{2,4} & 0 \\ 0 & R_{3,1} & R_{3,2} & R_{3,3} & R_{3,4} & 0 \\ 0 & R_{4,1} & R_{4,2} & R_{4,3} & R_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

\mathcal{I}_{green}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & G_{1,1} & G_{1,2} & G_{1,3} & G_{1,4} & 0 \\ 0 & G_{2,1} & G_{2,2} & G_{2,3} & G_{2,4} & 0 \\ 0 & G_{3,1} & G_{3,2} & G_{3,3} & G_{3,4} & 0 \\ 0 & G_{4,1} & G_{4,2} & G_{4,3} & G_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

\mathcal{I}_{blue}

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & B_{1,1} & B_{1,2} & B_{1,3} & B_{1,4} & 0 \\ 0 & B_{2,1} & B_{2,2} & B_{2,3} & B_{2,4} & 0 \\ 0 & B_{3,1} & B_{3,2} & B_{3,3} & B_{3,4} & 0 \\ 0 & B_{4,1} & B_{4,2} & B_{4,3} & B_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

$Kernel[:, :, 0]$

$$\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix}$$

$Kernel[:, :, 1]$

$$\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}$$

$Kernel[:, :, 2]$

$$\begin{bmatrix} w_{19} & w_{20} & w_{21} \\ w_{22} & w_{23} & w_{24} \\ w_{25} & w_{26} & w_{27} \end{bmatrix}$$

$w = np.hstack((Kernel[:, :, 0].flatten(), Kernel[:, :, 1].flatten(), Kernel[:, :, 2].flatten()))$

$x = np.hstack((Window_{red}.flatten(), Window_{green}.flatten(), Window_{blue}.flatten()))$

$FeatureMap[0, 0, 0] = ReLU(w.dot(x) + bias) = ReLU(w^T x + bias)$

CNN: VGG-16 as example

2D Convolution : only 3x3 kernel, stride 1, zero padding of thickness 1



\mathcal{I}_{red}

$$\begin{pmatrix} 0 & \mathbf{0} & \mathbf{0} & \mathbf{0} & 0 & 0 \\ 0 & \mathbf{R_{1,1}} & \mathbf{R_{1,2}} & \mathbf{R_{1,3}} & R_{1,4} & 0 \\ 0 & \mathbf{R_{2,1}} & \mathbf{R_{2,2}} & \mathbf{R_{2,3}} & R_{2,4} & 0 \\ 0 & R_{3,1} & R_{3,2} & R_{3,3} & R_{3,4} & 0 \\ 0 & R_{4,1} & R_{4,2} & R_{4,3} & R_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

\mathcal{I}_{green}

$$\begin{pmatrix} 0 & \mathbf{0} & \mathbf{0} & \mathbf{0} & 0 & 0 \\ 0 & \mathbf{G_{1,1}} & \mathbf{G_{1,2}} & \mathbf{G_{1,3}} & G_{1,4} & 0 \\ 0 & \mathbf{G_{2,1}} & \mathbf{G_{2,2}} & \mathbf{G_{2,3}} & G_{2,4} & 0 \\ 0 & G_{3,1} & G_{3,2} & G_{3,3} & G_{3,4} & 0 \\ 0 & G_{4,1} & G_{4,2} & G_{4,3} & G_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

\mathcal{I}_{blue}

$$\begin{pmatrix} 0 & \mathbf{0} & \mathbf{0} & \mathbf{0} & 0 & 0 \\ 0 & \mathbf{B_{1,1}} & \mathbf{B_{1,2}} & \mathbf{B_{1,3}} & B_{1,4} & 0 \\ 0 & \mathbf{B_{2,1}} & \mathbf{B_{2,2}} & \mathbf{B_{2,3}} & B_{2,4} & 0 \\ 0 & B_{3,1} & B_{3,2} & B_{3,3} & B_{3,4} & 0 \\ 0 & B_{4,1} & B_{4,2} & B_{4,3} & B_{4,4} & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

$Kernel[:, :, 0]$

$$\begin{bmatrix} w_1 & w_2 & w_3 \\ w_4 & w_5 & w_6 \\ w_7 & w_8 & w_9 \end{bmatrix}$$

$Kernel[:, :, 1]$

$$\begin{bmatrix} w_{10} & w_{11} & w_{12} \\ w_{13} & w_{14} & w_{15} \\ w_{16} & w_{17} & w_{18} \end{bmatrix}$$

$Kernel[:, :, 2]$

$$\begin{bmatrix} w_{19} & w_{20} & w_{21} \\ w_{22} & w_{23} & w_{24} \\ w_{25} & w_{26} & w_{27} \end{bmatrix}$$

$w = np.hstack((Kernel[:, :, 0].flatten(), Kernel[:, :, 1].flatten(), Kernel[:, :, 2].flatten()))$

$x = np.hstack((\textcolor{red}{Window}.flatten(), \textcolor{green}{Window}.flatten(), \textcolor{blue}{Window}.flatten()))$

$FeatureMap[1, 0, 0] = ReLU(w.dot(x) + bias) = ReLU(w^T x + bias)$

CNN: VGG-16 as example



Feature maps & Max pooling

- as many feature maps as there are kernels
- each kernel is detecting a particular feature (edges, texture,...)
- by adding more kernels, the model can learn to detect more complex features
- max pooling \Rightarrow shrinking of the feature maps
 - no learnable parameters
 - 2x2 kernel
 - stride 2
 - no zero padding

$$\begin{pmatrix} 1 & 1 & 2 & 4 \\ 5 & 6 & 7 & 8 \\ 3 & 2 & 1 & 0 \\ 1 & 2 & 3 & 4 \end{pmatrix} \xrightarrow{\text{max pooling}} \begin{pmatrix} 6 & 8 \\ 3 & 4 \end{pmatrix}$$

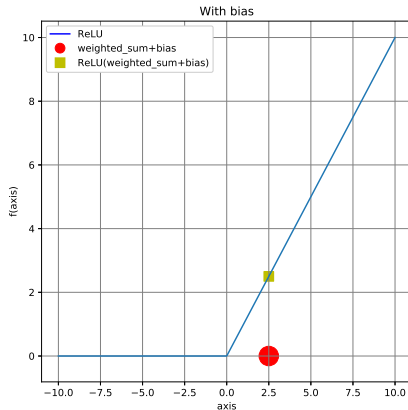
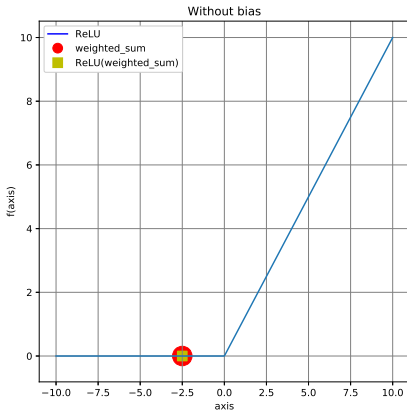
CNN: VGG-16 as example

Bias



$$weighted_sum = \sum_{i=1}^n w_i x_i = w^T x$$

- biases are learned parameters
- each neuron has a bias



Plan

- 1 Introduction
- 2 CNN
- 3 Xception**
- 4 VGG-16
- 5 ResNet-50
- 6 Conclusion
- 7 References

71 trainable layers ; 22,910,480 parameters



Layer	Type	Activation function	Output Shape	Param
input_1	(InputLayer)	N/A	(299, 299, 3)	0
block1_conv1	(Conv2D)	N/A	(149, 149, 32)	864
block1_conv1_bn	(Batch Normalization)	N/A	(149, 149, 32)	128
block1_conv1_act	(Activation)	ReLU	(149, 149, 32)	0
block1_conv2	(Conv2D)	N/A	(147, 147, 64)	18,432
block1_conv2_bn	(Batch Normalization)	N/A	(147, 147, 64)	256
block1_conv2_act	(Activation)	ReLU	(147, 147, 64)	0
block2...	(...)	...	(147, 147, 128)	...
conv2d_45	(Conv2D)	N/A	(74,74,128)	8192
block2_pool	(MaxPooling2D)	N/A	(74,74,128)	0
bn_45	(Batch Normalization)	N/A	(74,74,128)	512
add	(Add)	N/A	(74,74,128)	0
block3...	(...)	...	(74, 74, 256)	...
...	(...)	...	(...)	...
block14...	(...)	...	(10, 10, 1536)	1,582,080
avg_pool	(GlobalAveragePooling2D)	N/A	(, 2048)	0
predictions	(Dense)	Softmax	(, 1000)	2,049,000

Xception

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(299, 299, 3)
	(Conv2D)	N/A	(149, 149, 32)
	(Batch Normalization)	N/A	(149, 149, 32)
	(Activation)	ReLU	(149, 149, 32)
	(...)	...	(...)
	(...)	...	(...)
train this layer	(GlobalAveragePooling2D)	N/A	(, 2048)
	(Dense)	ReLU	
	(Dense)	Softmax	(, 1000)
	(Dense)	Softmax	(, nbClasses)

Trainable params: 6,147

Xception

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(299, 299, 3)
	(Conv2D)	N/A	(149, 149, 32)
	(Batch Normalization)	N/A	(149, 149, 32)
	(Activation)	ReLU	(149, 149, 32)
	(...)	...	(...)
	(...)	...	(...)
	(GlobalAveragePooling2D)	N/A	(, 2048)
train this layer	(Dense)	ReLU	(, 1000)
	(Dense)	Softmax	(, 1000)
	(Dropout)	N/A	(, 2048)
	(Dense)	Softmax	(, nbClasses)

Trainable params: 6,147

Plan

- 1 Introduction
- 2 CNN
- 3 Xception
- 4 VGG-16**
- 5 ResNet-50
- 6 Conclusion
- 7 References

VGG-16

Architecture (2014)

16 trainable layers ; 138,357,544 parameters



Layer	Type	Activation function	Output Shape	Param
input_1	(InputLayer)	N/A	(224, 224, 3)	0
block1_conv1	(Conv2D)	ReLU	(224, 224, 64)	1,792
block1_conv2	(Conv2D)	ReLU	(224, 224, 64)	36,928
block1_pool	(MaxPooling2D)	N/A	(112, 112, 64)	0
block2_conv1	(Conv2D)	ReLU	(112, 112, 128)	73,856
block2_conv2	(Conv2D)	ReLU	(112, 112, 128)	147,584
block2_pool	(MaxPooling2D)	N/A	(56, 56, 128)	0
block3_conv1	(Conv2D)	ReLU	(56, 56, 256)	295,168
block3_conv2	(Conv2D)	ReLU	(56, 56, 256)	590,080
block3_conv3	(Conv2D)	ReLU	(56, 56, 256)	590,080
block3_pool	(MaxPooling2D)	N/A	(28, 28, 256)	0
block4_conv1	(Conv2D)	ReLU	(28, 28, 512)	118,0160
block4_conv2	(Conv2D)	ReLU	(28, 28, 512)	2,359,808
block4_conv3	(Conv2D)	ReLU	(28, 28, 512)	2,35,9808
block4_pool	(MaxPooling2D)	N/A	(14, 14, 512)	0
block5_conv1	(Conv2D)	ReLU	(14, 14, 512)	2,359,808
block5_conv2	(Conv2D)	ReLU	(14, 14, 512)	2,359,808
block5_conv3	(Conv2D)	ReLU	(14, 14, 512)	2,359,808
block5_pool	(MaxPooling2D)	N/A	(7, 7, 512)	0
flatten	(Flatten)	N/A	(, 25088)	0
fc1	(Dense)	ReLU	(, 4096)	102,764,544
fc2	(Dense)	ReLU	(, 4096)	16,781,312
predictions	(Dense)	Softmax	(, 1000)	4,097,000

VGG-16

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(224, 224, 3)
	(Conv2D)	ReLU	(224, 224, 64)
	(Conv2D)	ReLU	(224, 224, 64)
	(MaxPooling2D)	N/A	(112, 112, 64)
	(Conv2D)	ReLU	(112, 112, 128)
	(Conv2D)	ReLU	(112, 112, 128)
	(MaxPooling2D)	N/A	(56, 56, 128)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(56, 56, 256)
	(MaxPooling2D)	N/A	(28, 28, 256)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(28, 28, 512)
	(MaxPooling2D)	N/A	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
train this layer	(MaxPooling2D)	N/A	(7, 7, 512)
	(Flatten)	N/A	(, 25088)
	(Dense)	ReLU	(, 4096)
	(Dense)	ReLU	(, 4096)
	(Dense)	Softmax	(, 1000)
	(Dense)	Softmax	(, nbClasses)

VGG-16

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(224, 224, 3)
	(Conv2D)	ReLU	(224, 224, 64)
	(Conv2D)	ReLU	(224, 224, 64)
	(MaxPooling2D)	N/A	(112, 112, 64)
	(Conv2D)	ReLU	(112, 112, 128)
	(Conv2D)	ReLU	(112, 112, 128)
	(MaxPooling2D)	N/A	(56, 56, 128)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(56, 56, 256)
	(MaxPooling2D)	N/A	(28, 28, 256)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(28, 28, 512)
	(MaxPooling2D)	N/A	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
	(Conv2D)	ReLU	(14, 14, 512)
train this layer	(MaxPooling2D)	N/A	(7, 7, 512)
	(Flatten)	N/A	(, 25088)
	(Dense)	ReLU	(, 4096)
	(Dense)	ReLU	(, 4096)
	(Dense)	Softmax	(, 1000)
	(Dropout)	N/A	(, 4096)
	(Dense)	Softmax	(, nbClasses)

Plan

- 1 Introduction
- 2 CNN
- 3 Xception
- 4 VGG-16
- 5 ResNet-50**
- 6 Conclusion
- 7 References

RESNET-50

Architecture (2015)

50 trainable layers ; 25,636,712 parameters



Layer	Type	Activation function	Output Shape	Param
input_1	(InputLayer)	N/A	(224, 224, 3)	0
res...branch..	(Conv2D)	N/A	(112, 112, 64)	9,472
bn...branch..	(Batch Normalization)	N/A	(112, 112, 64)	256
activation...	(Activation)	ReLU	(112, 112, 64)	0
max_pooling2d_1	(MaxPooling2D)	N/A	(56, 56, 64)	0
res...branch..	(Conv2D)	N/A	(56, 56, 64)	4,160
bn...branch..	(Batch Normalization)	N/A	(56, 56, 64)	256
activation...	(Activation)	ReLU	(56, 56, 64)	0
res...branch..	(Conv2D)	N/A	(56, 56, 64)	36,928
bn...branch..	(Batch Normalization)	N/A	(56, 56, 64)	256
activation...	(Activation)	ReLU	(56, 56, 64)	0
res...branch..	(Conv2D)	N/A	(56, 56, 256)	16,640
bn...branch..	(Batch Normalization)	N/A	(56, 56, 256)	256
activation...	(Activation)	ReLU	(56, 56, 64)	0
res...branch..	(Conv2D)	ReLU	(28, 28, 128)	32,896
res...branch..	(Conv2D)	ReLU	(28, 28, 128)	147,584
res...branch..	(Conv2D)	ReLU	(28, 28, 512)	66,048
res...branch..	(Conv2D)	ReLU	(14, 14, 256)	131,328
res...branch..	(Conv2D)	ReLU	(14, 14, 256)	590,080
res...branch..	(Conv2D)	ReLU	(14, 14, 1024)	263,168
res...branch..	(Conv2D)	ReLU	(7, 7, 512)	524,800
res...branch..	(Conv2D)	ReLU	(7, 7, 512)	2,359,808
res...branch..	(Conv2D)	ReLU	(7, 7, 2048)	1,050,624
avg_pool	(GlobalAveragePooling2D)	N/A	(,2048)	0
fc1000	(Dense)	SoftMax	(, 1000)	2,049,000

RESNET-50

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(224,224,3)
	(Conv2D)	ReLU	(112, 112, 64)
	(MaxPooling2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(28, 28, 128)
	(Conv2D)	ReLU	(28, 28, 128)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(14, 14, 256)
	(Conv2D)	ReLU	(14, 14, 256)
	(Conv2D)	ReLU	(14, 14, 1024)
	(Conv2D)	ReLU	(7, 7, 512)
	(Conv2D)	ReLU	(7, 7, 512)
	(Conv2D)	ReLU	(7, 7, 2048)
	(GlobalAveragePooling2D)	ReLU	(,2048)
train this layer	(Dense)	Softmax	(,1000)
	(Dense)	Softmax	(, nbClasses)

Trainable params: 6,147

RESNET-50

Transfer Learning



	Layer Type	Activation function	Output Shape
freeze weights learned on ImageNet	(InputLayer)	N/A	(224,224,3)
	(Conv2D)	ReLU	(112, 112, 64)
	(MaxPooling2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 64)
	(Conv2D)	ReLU	(56, 56, 256)
	(Conv2D)	ReLU	(28, 28, 128)
	(Conv2D)	ReLU	(28, 28, 128)
	(Conv2D)	ReLU	(28, 28, 512)
	(Conv2D)	ReLU	(14, 14, 256)
	(Conv2D)	ReLU	(14, 14, 256)
	(Conv2D)	ReLU	(14, 14, 1024)
	(Conv2D)	ReLU	(7, 7, 512)
	(Conv2D)	ReLU	(7, 7, 512)
	(Conv2D)	ReLU	(7, 7, 2048)
train this layer	(GlobalAveragePooling2D)	ReLU	(,2048)
	(Dense)	Softmax	(,1000)
	(Dropout)	N/A	(, 2048)
	(Dense)	Softmax	(, nbClasses)

Trainable params: 6,147

Plan

- 1 Introduction
- 2 CNN
- 3 Xception
- 4 VGG-16
- 5 ResNet-50
- 6 Conclusion**
- 7 References

Conclusion

Our results



Yann LeCun

@ylecun

Follow



Training with large minibatches is bad for your health.
More importantly, it's bad for your test error.
Friends dont let friends use minibatches larger than 32. arxiv.org/abs/1804.07612

2:00 PM - 26 Apr 2018

Conclusion

Our results



Training set (3867 images)

Validation set (967 images) : Top 1 Accuracy												
With 0.5 dropout							Without dropout					
Epochs	2			10			2			10		
Batch size	8	16	32	8	16	32	8	16	32	8	16	32
Xception	0.3588	0.3661	0.3609	0.3909	0.3433	0.3454	0.4116	0.4012	0.4012	0.4199	0.3899	0.3733
VGG-16	0.9741	0.9835	0.9824	0.9824	0.9855	0.9824	0.9762	0.9845	0.9866	0.9814	0.9814	0.9814
ResNet-50	0.9659	0.9700	0.9721	0.9731	0.9741	0.9731	0.9690	0.9690	0.9659	0.9710	0.9710	0.9710

Test set (1209 images) : Top 1 Accuracy												
With 0.5 dropout							Without dropout					
Epochs	2			10			2			10		
Batch size	8	16	32	8	16	32	8	16	32	8	16	32
Xception	0.3524	0.3677	0.3490	0.4030	0.3711	0.3708	0.3984	0.3764	0.3423	0.3667	0.3422	0.3598
VGG-16	0.9586	0.9702	0.9727	0.9644	0.9702	0.9744	0.9628	0.9686	0.9653	0.9669	0.9694	0.9686
ResNet-50	0.9661	0.9661	0.9711	0.9702	0.9735	0.9711	0.9686	0.9694	0.9694	0.9537	0.9644	0.9639

Plan

- 1 Introduction
- 2 CNN
- 3 Xception
- 4 VGG-16
- 5 ResNet-50
- 6 Conclusion
- 7 References**

References

- Stanford University School of Engineering
 - Lecture 5 | Convolutional Neural Networks
<https://www.youtube.com/watch?v=bNb2fEVKeEo>
 - Lecture 7 | Training Neural Networks II
https://www.youtube.com/watch?v=_JB0A07QxSA (1h12min : transfer learning)
 - Lecture 9 | CNN Architectures
<https://www.youtube.com/watch?v=DA0cjicFr1Y>
 - CS231n Convolutional Neural Networks for Visual Recognition
<http://cs231n.github.io/convolutional-networks/>
- MIT OpenCourseWare
 - 12a: Neural Nets
<https://www.youtube.com/watch?v=uXt8qF2Zzfo> bias: 23min
 - 12b: Deep Neural Nets
https://www.youtube.com/watch?v=VrMHA3yX_QI (kernels = neurone: 12min; softmax explanation: 34min)

References

- AlexNet Original Article:
<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>
- VGG-16 Original article : VERY DEEP CONVOLUTIONAL NETWORKS FOR LARGE-SCALE IMAGE RECOGNITION
<https://arxiv.org/pdf/1409.1556.pdf>
- VGG-16: "The only preprocessing we do is subtracting the mean RGB value, computed on the training set, from each pixel."
<https://machinelearningmastery.com/use-pre-trained-vgg-model-classify-objects-photographs/>)
- Batch Normalization : learnable parameters
<https://medium.com/coinmonks/paper-review-of-alexnet-caffenet-winner-in-ilsvrc-2012-image-classification-b>
- Keras pretrained Models
<https://keras.io/applications/>
- Pretrained Deep Neural Networks Matlab
- <https://www.mathworks.com/help/deeplearning/ug/pretrained-convolutional-neural-networks.html>